Docker is great for containerization, but when you scale applications, you hit limitations:

Drawbacks of Docker (Loopholes in Using Just Docker)

1.Single Host Limitation

- 1. Docker runs containers on a single machine.
- 2. What if traffic increases? We need multiple machines.
- 3. Solution? Kubernetes (K8s) helps in scaling across multiple nodes.

2. Manual Scaling

- 1. In Docker, you **manually start** and **stop** containers.
- 2. No auto-scaling, no self-healing.
- 3. Kubernetes can auto-scale and restart failed containers.

3.No Load Balancing

- 1. Docker itself doesn't manage traffic distribution.
- 2. If multiple containers are running, how will traffic be distributed?
- 3. Kubernetes has built-in load balancing.

4. Storage & Networking Challenges

- 1. Docker doesn't handle **persistent storage** properly.
- 2. Networking between multiple Docker hosts is **not easy**.
- 3. Kubernetes has Persistent Volumes (PV), Persistent Volume Claims (PVC), and Service Networking.

5.No High Availability (HA)

- 1. If a container crashes in Docker, it stays down unless restarted manually.
- 2. In Kubernetes, Pods auto-restart and maintain high availability.

Kubernetes vs. Docker-Compose:

Feature	Docker-Compose	Kubernetes
Scaling	Manual	Auto-scaled Pods
Load Balancing	No	Built-in Services & Ingress
Self-healing	No	Yes
Multi-node Support	No	Yes
Storage	Local	Persistent Volumes


Key Reasons to Use Kubernetes:


- 🔟 Scalability 🌠
- 2 Self-Healing & High Availability 🔷
- Load Balancing & Service Discovery
- 4 Multi-Cloud & Hybrid Deployments 🗅
- 5 Networking & Security CNI & RBAC
- 6 Persistent Storage

- Handles scaling automatically
- Recovers from failures automatically
- **Mate :** Distributes traffic intelligently
- Works across multiple clouds
- Ensures persistent storage & networking

1 EKS and 1 EC2 we can run 20 Applications, with high availability and Auto scaling, Loadbalencing. Backup u need according to

Your requirement, its your headache!!





RDS

connection strings inside pods

Mounted

volumes via PVC

Deployment Type	Resources Used	Max Applications	Cost per Hour	Monthly Cost (730 hrs)
EKS (1 EC2 as Worker Node)	1x t3.medium (2 vCPU, 4GB RAM)	10 Applications	\$0.1418 (EKS + EC2)	\$103.51


KUBERNETES ARCHITETURE

Master Node components:

API SERVER
KUBE SCHEDULAR
KUBE CONTROLLER MANAGER

Worker node Components:

Kubelet Kube-proxy Container runtime(CRI-O)



KUBERNETES CLUSTER

X How Kubernetes Works (Bird's Eye View):

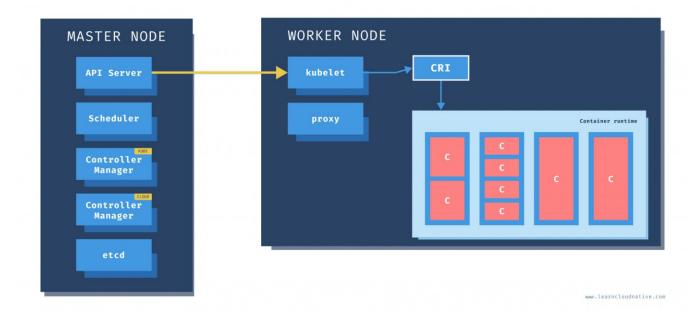
Kubernetes = **Cluster** of machines.

It has 2 major parts:

- 1.Master Node (Control Plane) brains <a>
- **2.Worker Nodes** workers (where your apps/pods actually run)

Master Node (a.k.a Control Plane Components):

This is where K8s makes all decisions like a boss 😇



Component	Role
API Server	riangle The front door! Every command hits this. kubectl $ o$ API server
etcd	
Controller Manager	Watches and reacts to changes (creates pods if one dies)
Scheduler	Decides where to run new pods (which node)
Cloud Controller Manager	Talks to cloud (AWS, GCP) if needed



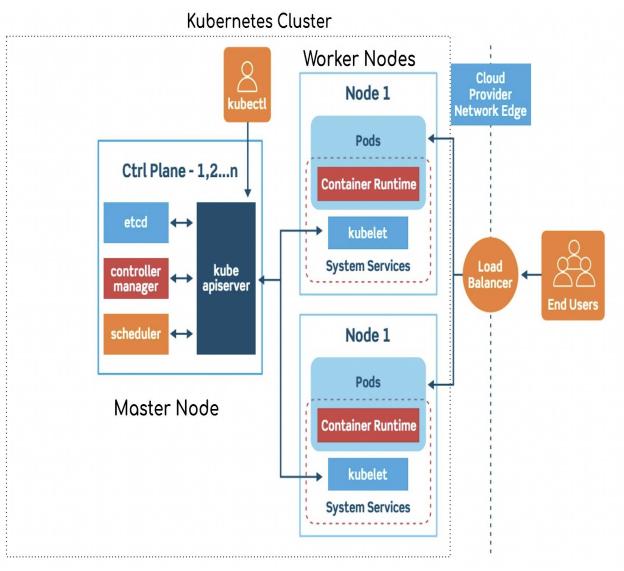
A **Node** in Kubernetes is a **worker machine (VM or physical server)** where the containers (Pods) run. Each **Kubernetes cluster**

- 1 Control Plane (Master Node) → Manages the cluster
- 2 Worker Nodes → Run workloads (containers inside Pods)

♦ Worker Node Components

Each worker node has these core components:

- 1 Kubelet (Agent)
- •The **bridge** between the Node and the Control Plane
- •Receives commands from the API Server
- •Ensures the Pod is running as expected
- Communicates with the Container Runtime
 - 2 Container Runtime (Docker/Containerd/CRI-O)
 - •Runs and manages containers inside the pod
 - •The default runtime for Kubernetes is containerd
 - **3** Kube Proxy (Networking Component)
 - •Handles networking inside the cluster
 - Enables Pod-to-Pod communication across nodes
 - •Implements Load Balancing for services



X Real-World Example:

If your master node dies:

- Existing pods may keep running (until they crash or need rescheduling)
- •But you can't deploy, scale, or recover anything → app becomes fragile.

If a worker node dies:

- Control plane detects it
- •Reschedules pods on other available nodes → self-healing kicks in

High Availability (HA):

- •Without a working master node, new pods can't be scheduled or managed.
- •Without worker nodes, there's **no place to run** the workloads.
- •So, both are mutually dependent to keep the app available and responsive.

Rule to write a Yaml:

X No tabs allowed – use only spaces

Indentation matters — usually 2 spaces per level

Key: value format Example: name: mypod

🗐 Lists use - (dash)

Strings don't need quotes (usually):

name: myapp # this is fine name: "myapp" # also valid

Boolean values must be lowercase:

enabled: true # 🛂

enabled: FALSE # 💢 (wrong)

Multiline strings:

note: | This is line 1 This is line 2

Use anchors (&) and aliases (*) for reuse:

default: &app defaults

image: nginx

port: 80

dev:

<<: *app_defaults

replicas: 2

Pods are the smallest deployable units of computing that you can create and manage in Kubernetes. A Pod (as in a pod of whales is a group of one or more containers, with shared storage and network resources, and a specification for how to run the containers.

pod.yaml:

apiVersion: v1

kind: Pod

metadata:

name: my-nginx-pod # Name of the Pod

labels:

app: nginx

spec:

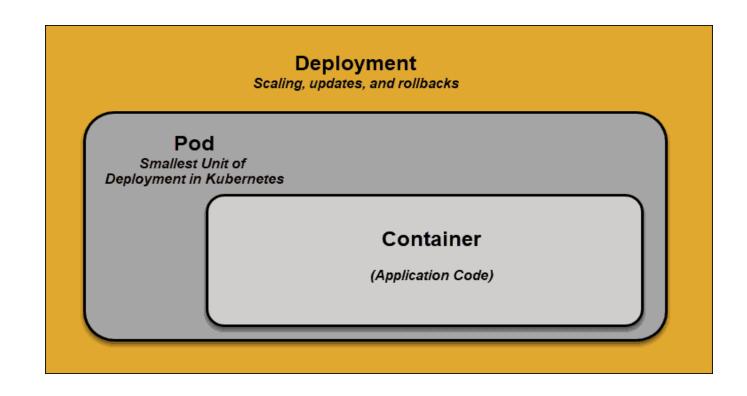
containers:

- name: nginx-container

image: nginx:latest

ports:

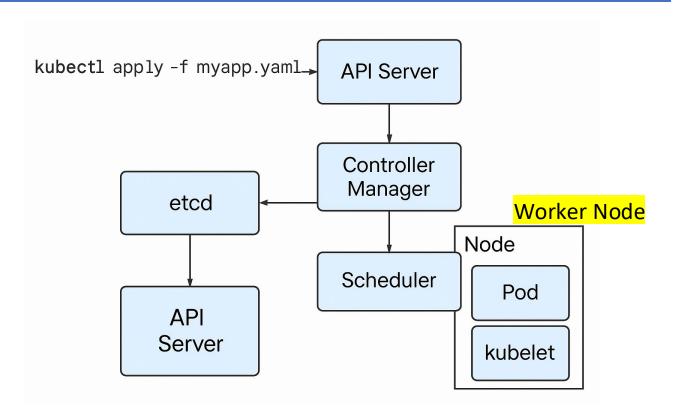
- containerPort: 80



- > kubectl apply -f pod.yaml
- > kubectl get pods

> kubectl describe pod my-nginx-pod > kubectl delete pod my-nginx-pod

- apiVersion: v1 → Using Kubernetes API version 1
- ♦ kind: Pod → Defines that this is a Pod resource
- metadata: → Assigns a name and labels to the pod
- containers: → List of containers inside the pod
- ♦ image: nginx:latest → Uses the official NGINX image from Docker Hub
- ◆ ports: → Exposes port 80 inside the container.
 - 1. How It Flows (Simple Steps):
 - 2. You type kubectl apply -f mypod.yaml
 - 3. kubectl \rightarrow hits API server
 - API server stores desired state in etcd
 - 5. Controller manager notices no pod is running
 - 6. Scheduler picks a node
 - 7. kubelet (on that node) pulls image, runs container
 - 8. 🤌 Pod is up!



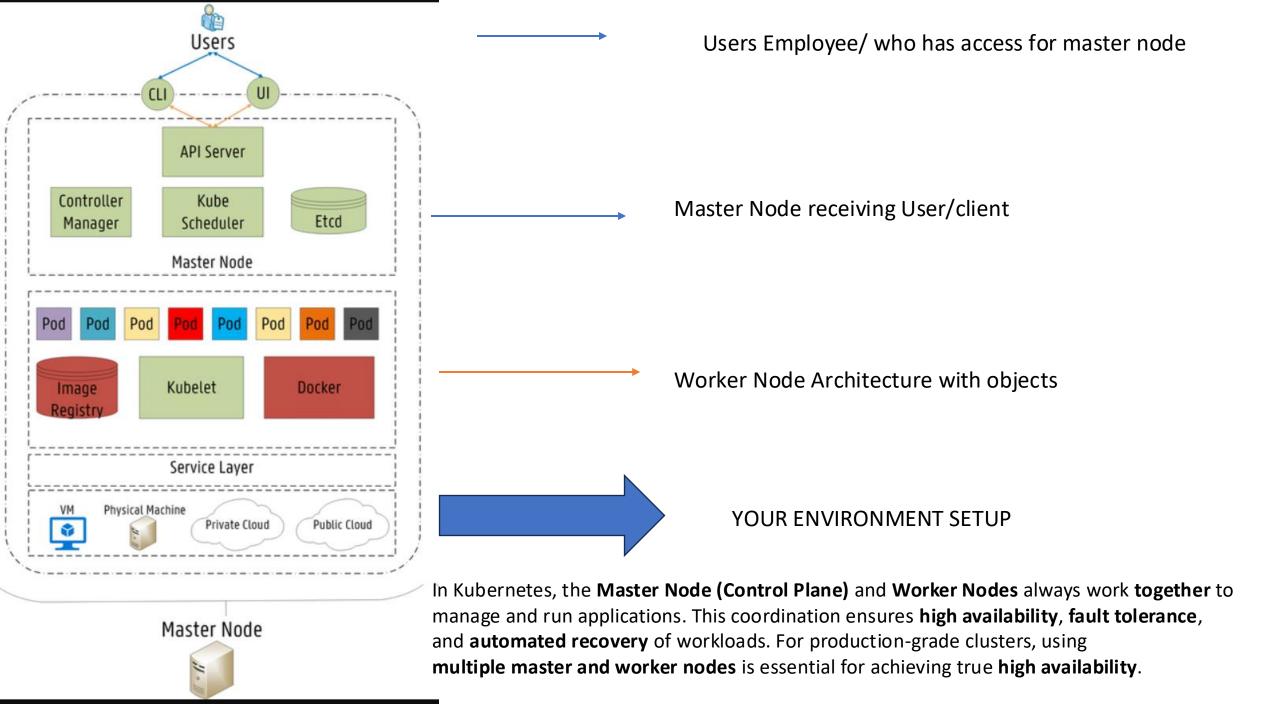


A **Pod** in Kubernetes goes through different states based on its lifecycle and health status. Below are the key **Pod health states**:

State	Description	
Pending	The Pod has been accepted by Kubernetes but not yet running (waiting for resources, scheduling, or pulling images).	
Running	The Pod is successfully scheduled and at least one container is in a Running state.	
Succeeded	All containers in the Pod have completed execution successfully (only applies to jobs/batch workloads).	
Failed	One or more containers in the Pod have terminated with a non-zero exit code.	
Unknown	The Pod state is unknown (often due to network or node communication issues).	

State	Description	
Waiting	The container is waiting to start. It might be pulling an image, waiting for resources, or experiencing an issue.	
Running	nning The container is running normally without issues.	
Terminated	The container has exited. It could be successful (Exit Code 0) or failed (Exit Code ≠ 0).	

• You may see Waiting with a reason, such as ImagePullBackOff, CrashLoopBackOff, or
ErrImagePull.



F	ew	commands to learn	8	Kubernetes Commands You Must
	#	Command		What It Does

kubectl get pods

kubectl get pods -A

kubectl get nodes

kubectl get deployments

kubectl describe pod <pod-

kubectl describe node < node-

kubectl logs <pod-name>

kubectl logs <pod> -c

kubectl exec -it <pod> --

kubectl apply -f app.yaml

kubectl create -f app.yaml

kubectl delete -f app.yaml

<container>

/bin/bash

kubectl get svc

kubectl get all

name>

name>

10

11

Kubernetes Commands You Must Know		
What It Does		
List all running pods in the current namespace		

15	kubectl delete pod <pod-name></pod-name>
16	kubectl scale deployment <deploy< td=""></deploy<>

Get pods in all namespaces

List all deployments

Show logs of a pod

List all nodes (workers + masters) in the cluster

Show all K8s objects: pods, svc, deployments, etc.

Show services (ClusterIP, NodePort, LB, etc.)

Detailed pod info + events + errors

Logs for a specific container in a pod

SSH into a running pod/container

Create/update K8s resources from YAML

Delete K8s resources defined in YAML

Create K8s resources from YAML (fails if already exists)

Detailed info of a specific node

kubectl scale deployment <deployname> --replicas=5

kubectl expose deployment <deploy-

kubectl edit deployment <deploy-name>

kubectl port-forward pod/<pod-name>

kubectl config use-context < context>

kubectl cp <pod-name>:<path> <local-

name> --type=NodePort --port=80

kubectl rollout status

kubectl rollout undo

deployment/<deploy-name>

deployment/<deploy-name>

Scale up/down pods

Expose app as a service

Check rollout progress

Rollback to previous deployment

Open live YAML editor for deployment

Access pod port locally (dev/testing)

Show pod-level resource usage

Show node-level resource usage

Switch between clusters (like

List PersistentVolumeClaims

Show client/server version

List Persistent Volumes

Show your current kubeconfig setup

Copy file from pod to local machine

Show schema/documentation for a

resource (great for YAML help!)

(CPU/RAM)

dev/stage/prod)

Delete a single pod

17

18

19

20

21

22

23

27

28

29

30

8080:80

kubectl top pod

kubectl top node

kubectl config view

kubectl get pvc

kubectl get pv

kubectl version --short

kubectl explain pod

path>

- 1 ReplicationController (RC) [Old]
- •Ensures a **specified number of pod replicas** are running at all times.
- •If a pod crashes, it creates a new one.
- •Obsolete → Replaced by ReplicaSet.
 - 2 ReplicaSet (RS) [Improved RC]
- •Works like ReplicationController but supports **label selectors with set-based selectors** (more flexible).
- Ensures the desired number of pod replicas are running.
- •If a pod fails, it replaces it with a new one.
- 3 Deployment [Higher-Level Controller]
- •Manages ReplicaSets automatically and provides additional features like:
- Rolling updates (zero downtime updates)
- Rollback to previous versions
- **Declarative updates** (modify pod templates easily)
- Self-healing & scaling

```
Replicaset.yaml:
apiVersion: apps/v1
kind: ReplicaSet
metadata:
 name: my-nginx-rs
 labels:
  app: nginx
spec:
 replicas: 3
 selector:
  matchLabels:
   app: nginx
 template:
  metadata:
   labels:
    app: nginx
  spec:
   containers:
   - name: nginx-container
    image: nginx:latest
    ports:
    - containerPort: 80
```

Why Prefer Deployment Over ReplicaSet?

Feature	ReplicaSet	Deployment
Ensures Pod Count	✓ Yes	✓ Yes
Rolling Updates	× No	✓ Yes
Rollbacks	× No	✓ Yes
Auto-Manages ReplicaSet	× No	✓ Yes
Recommended Usage	X No (Only for special cases)	Yes (Preferred)

Use Deployments for managing apps, not ReplicaSets directly.

kubectl apply -f deployment.yaml kubectl get deployments

Rollout: kubectl rollout undo deployment my-nginx-deployment

Feature	Description	
☑ Declarative Control	You tell K8s what state you want, it makes it happen.	
Rolling Updates	Update app without downtime. K8s replaces old pods slowly.	
& Self-Healing	If a pod crashes, Deployment brings it back automatically.	
	Easy to increase/decrease pod count with 1 command.	
Version Control	Roll back to a previous version if something breaks.	
Safe & Controlled	Can pause/resume/monitor updates — great for production apps.	

deployment.yaml:

apiVersion: apps/v1 kind: Deployment metadata:

name: my-nginx-deployment

labels:

app: nginx

spec:

replicas: 3

selector:

matchLabels:

app: nginx

template:

metadata:

labels:

app: nginx

spec:

containers:

- name: nginx-container

image: nginx:latest

ports:

- containerPort: 80

♦ Labels & Selectors (For Grouping and Identifying Resources)

Labels are key-value pairs attached to Kubernetes objects (Pods, Nodes, etc.).

Selectors help filter and select resources based on labels.

NOTE: When we create multiple Pods using a Deployment, labels are defined inside the Pod template.

The Deployment uses selector.matchLabels to find and manage the Pods based on those labels.

Taints & Tolerations (For Node-Pod Restrictions)

Taints prevent certain Pods from running on a Node unless the Pod has a matching

Toleration.

f Example: Taint a Node

kubectl taint nodes my-node key=value:NoSchedule

apiVersion: v1

kind: Pod

metadata:

name: my-tolerated-pod

spec:

tolerations:

- key: "key"

operator: "Equal"

value: "value"

effect: "NoSchedule"

containers:

- name: nginx

image: nginx:latest

apiVersion: v1

kind: Pod

metadata:

name: my-nginx-pod

labels:

app: nginx

env: production

spec:

containers:

- name: nginx

image: nginx:latest

selector:

matchLabels:

app: nginx



You define labels under template.metadata.labels (inside the Deployment) selector.matchLabels in the Deployment spec looks for those labels Kubernetes then manages only the Pods that match those labels



Node Affinity allows you to **control which nodes a pod can be scheduled on** based on node labels.

It's an advanced version of nodeSelector that supports:

Soft (preferred) rules \rightarrow Pod *tries* to be scheduled on preferred nodes.

Hard (required) rules \rightarrow Pod *must* be scheduled on specified nodes.

Types of Node Affinity

Туре	Effect
requiredDuringSchedulingIgnoredDuringExecution	Hard rule \rightarrow Pod must be scheduled on the specified node.
preferred During Scheduling Ignored During Execution	Soft rule → Scheduler tries to place the Pod on preferred nodes but will schedule elsewhere if needed.

How is Node Affinity Different from Taints & Tolerations?

Feature	Node Affinity	Taints & Tolerations
Who sets it?	Pod definition	Node definition
Purpose	Control pod placement based on labels	Restrict certain pods from running on nodes
Effect	Scheduler decides based on preferences	Node blocks pods unless they have tolerations

nodeaffinity.yaml:

apiVersion: v1

kind: Pod metadata:

name: my-app-pod

spec:

affinity:

nodeAffinity:

required During Scheduling Ignored During Execution:

nodeSelectorTerms:

- matchExpressions:

- key: env

operator: In

values:

- production

containers:

- name: my-app

image: nginx



Requests and Limits in Kubernetes



What are Requests and Limits?

Requests and Limits control resource allocation (CPU & Memory) for Pods in Kubernetes.

- •Requests → Minimum resources a container gets.
- •Limits → Maximum resources a container can use.
- Why use them?
- Prevents a single container from consuming all resources.
- Ensures fair resource distribution across Pods.
- Helps Kubernetes **schedule** Pods efficiently.

- What Happens If a Pod Exceeds Its Limits?
- If CPU exceeds the limit → Pod is throttled (slows down but doesn't crash).
- ✓ If Memory exceeds the limit → Pod is OOMKilled (Out of Memory).

requestandlimits,.yaml:

```
apiVersion: v1
kind: Pod
metadata:
 name: my-app
spec:
 containers:
 - name: nginx
  image: nginx
  resources:
   requests:
    memory: "128Mi" # Minimum memory
                  # Minimum CPU (0.25 cores)
    cpu: "250m"
   limits:
    memory: "256Mi" # Max memory allowed
    cpu: "500m"
                  # Max CPU (0.5 cores)
```



Pod gets at least 128Mi memory and 0.25 CPU

Cannot exceed 256Mi memory and 0.5 CPU



Pod Autoscaling in Kubernetes



What is HPA?

HPA (Horizontal Pod Autoscaler) automatically scales the number of pods in a deployment based on CPU, memory, or custom metrics.

- - Why use HPA?
- Handles traffic spikes automatically.
- Prevents unnecessary over-provisioning.
- Optimizes resource utilization and costs.
- - **How HPA Works**
 - Monitors **CPU/Memory usage** of running pods.
- If usage exceeds the defined threshold, HPA adds more pods.
- If usage drops below the threshold, HPA removes extra pods.
- Uses metrics-server to get resource usage.

kubectl top pods

kubectl get hpa

kubectl apply -f https://github.com/kubernetes-sigs/metricsserver/releases/latest/download/components.yaml

autoscaling.yaml:

apiVersion: autoscaling/v2

kind: HorizontalPodAutoscaler

metadata:

name: my-app-hpa

spec:

scaleTargetRef:

apiVersion: apps/v1

kind: Deployment

name: my-app-deployment

minReplicas: 2

maxReplicas: 10

metrics:

- type: Resource

resource:

name: cpu

target:

type: Utilization

averageUtilization: 50 # Scale if CPU > 50%

HPA vs VPA vs Cluster Autoscaler

Feature	HPA (Horizontal)	VPA (Vertical)	Cluster Autoscaler
Scaling Type	More Pods	Bigger Pods	More Nodes
Triggers	CPU, Memory, Custom Metrics	CPU, Memory	Node resource needs
Use Case	Handle high traffic	Optimize per-pod resources	Scale entire cluster

How Metrics Server Works in Kubernetes

****** What is Metrics Server?

The Metrics Server is a Kubernetes component that collects resource usage metrics (CPU & Memory) from nodes and pods and provides them to HPA (Horizontal Pod Autoscaler), VPA (Vertical Pod Autoscaler), and kubectl top commands.

P Why do we need it?

Enables HPA & VPA for auto-scaling.

Provides real-time CPU & Memory metrics of Pods and Nodes.

Lightweight & efficient, unlike full monitoring tools like Prometheus.

kubectl get pods -n kube-system | grep metrics-server

Metrics Server vs Prometheus

Feature	Metrics Server	Prometheus
Purpose	Provides CPU & memory usage	Full monitoring & alerting
Use Case	HPA, kubectl top	Detailed metrics & dashboards
Data Storage	Does not store data	Stores historical metrics
Complexity	Simple & lightweight	More complex setup

Health Probes (Liveness, Readiness, and Startup Probes)

Kubernetes uses **probes** to check if a container is **healthy** and **ready** to serve traffic.

Probe Type	Purpose
Liveness Probe	Checks if the container is alive. If it fails, Kubernetes restarts the container.
Readiness Probe	Checks if the container is ready to accept traffic. If it fails, traffic is not sent to the pod.
Startup Probe	Ensures the app fully starts before checking liveness.

- If the liveness probe fails, the pod is restarted.
 - If the readiness probe fails, the pod does not receive traffic.
- > Probes are used to check the health of containers in a Pod.
- > There are three types: Liveness, Readiness, and Startup probes.
- > Liveness Probe checks if the app is still running. If not, the container is restarted.
- ➤ Readiness Probe checks if the app is ready to serve traffic. If not, traffic is stopped.
- ➤ Startup Probe is used for apps that take time to start it delays other probes until it's ready.
- Probes can work using httpGet, exec, or tcpSocket.
- > They help in self-healing, high availability, and zero-downtime deployments.
- > Without probes, Kubernetes won't know if the app is stuck or unhealthy.

myprobe.yaml:

```
apiVersion: v1
kind: Pod
metadata:
 name: my-app
spec:
 containers:
 - name: my-container
  image: nginx
  livenessProbe:
   httpGet:
    path: /
    port: 80
   initialDelaySeconds: 5
   periodSeconds: 10
  readinessProbe:
   httpGet:
    path: /
    port: 80
   initialDelaySeconds: 3
   periodSeconds: 5
```

ConfigMaps (Manage Configuration in Kubernetes):

ConfigMaps store configuration data (key-value pairs) separately from the application code.



Keeps configuration separate from application code.

Allows dynamic updates without rebuilding the container.

Zan be used in environment variables, command arguments, or mounted as files.

How to Use the ConfigMap in a Pod:

ConfigMap as a file inside the container

ConfigMap.yaml

apiVersion: v1 kind: ConfigMap

metadata:

name: app-config

data:

database_url: "mongodb://localhost:27017"

app_mode: "production"

log level: "info"

apiVersion: v1

kind: Pod

metadata:

name: app-pod

spec:

containers:

- name: app-container

image: myapp:latest

envFrom:

- configMapRef:

name: app-config

apiVersion: v1

kind: Pod

metadata:

name: app-pod

spec:

containers:

- name: app-container

image: myapp:latest

volumeMounts:

name: config-volume mountPath: /etc/config

volumes:

- name: config-volume

configMap:

name: app-config

Namespaces (Isolate Resources in Kubernetes)

Namespaces **logically separate Kubernetes resources** in the same cluster.

echo -n "admin" | base64 # YWRtaW4=

echo -n "password123" | base64 # cGFzc3dvcmQxMjM=

Why use Namespaces?

Helps **organize** resources in large clusters.

Enables RBAC (Role-Based Access Control) per namespace.

Encoding base64

Supports **resource quotas** per team/project.

Namespace.yaml

apiVersion: v1

kind: Namespace

metadata:

name: dev-environment

Secrets.yaml

apiVersion: v1

kind: Secret

metadata:

name: my-secret

type: Opaque

data:

username: YWRtaW4= # Base64 "admin"

password: cGFzc3dvcmQxMjM= # Base64 encoded "

Declarative creating namespace:

Kubeclt create ns <ns name>

Secrets in Kubernetes:



What is a Secret?

A **Secret** is a Kubernetes object used to **store sensitive data** such as passwords, API keys, or TLS certificates securely.

Why use Secrets instead of ConfigMaps?

Encoded (Base64) data for security.

Prevents storing secrets in container images.

Can be mounted as files or environment variables.

What is a Service in Kubernetes?

A Service is a Kubernetes resource that exposes a set of Pods as a network-accessible endpoint.

Since Pods are **ephemeral** (they can be restarted, deleted, or rescheduled), a Service ensures **stable communication** between them.

Why Use Services?

Provides a stable IP address & DNS name.

Load-balances traffic to multiple Pods.

Allows external access to internal applications.

Enables communication between different components in a cluster.

Types of Kubernetes Services

Service Type	Use Case	Accessible From	Example
ClusterIP (default)	Internal communication	Inside the cluster only	Microservices
NodePort	Expose service on a fixed port	Any node's IP + port	Basic external access
LoadBalancer	Publicly expose service	Internet via cloud provider	Web applications
ExternalName	Maps service to an external domain	External (DNS resolution only)	Connecting to external DBs

ClusterIP.yaml

apiVersion: v1 kind: Service metadata:

name: my-clusterip-service

spec:

selector:

app: my-app

ports:

- protocol: TCP

port: 80 # Service port

targetPort: 8080 # Pod's container

port

type: ClusterIP

LoadBalancer

apiVersion: v1 kind: Service metadata:

name: loadbalancer-service

spec:

selector:

app: myapp

ports:

- protocol: TCP

port: 80

targetPort: 8080

type: LoadBalancer

NodePort Service

apiVersion: v1 kind: Service metadata:

name: nodeport-service

spec:

selector:

app: myapp

ports:

- protocol: TCP

port: 80

targetPort: 8080

nodePort: 30000

type: NodePort

A **ClusterIP** service exposes the application on an internal IP within the Kubernetes cluster.

A LoadBalancer service exposes the application externally and can be used with cloud providers to allocate an external IP.

A **NodePort** service exposes the application on a static port on each node's IP, which can be accessed externally.



Understanding Pod Disruption Budgets (PDB) in Kubernetes!

In a highly available microservices environment, maintaining application stability during voluntary disruptions (like node upgrades, scaling, or maintenance) is critical. This is where Pod Disruption Budgets (PDBs) come into play!



◆ What is a PDB?

A Pod Disruption Budget ensures that a minimum number of pods remain available during voluntary disruptions. It helps maintain service availability when nodes are drained or updated.

ndh vaml



◆ Key PDB Parameters:

ula ana Ni ani ana/ilalin a Ni anna ina maliana



minAvailable → The minimum number of pods that must be running.

maxUnavailable -> The maximum number of pods that can be disrupted at a time

when Nodes(blue) are cord	on:		pub.yami.	
PDB_	Deployment:	pod status:	apiVersion: policy/v	
minAvailable: 3	replicas: 5	3 running 2 pending	metadata: name: pdb-minav	ailable-count
minAvailable: 1	replicas: 5	1 Running 4 pending	spec: minAvailable: 2	// maxUnavailable: 1
maxUnavilable: 2	replicas: 5	3 running 2 pending	selector: matchLabels: app: myapp	

Command NODE/POD Level	Purpose
kubectl exec -it <pod-name> /bin/sh</pod-name>	Get into container shell (fix config or test app manually)
kubectl debug node/ <node-name>image=busybox</node-name>	Debug node with ephemeral container
kubectl describe events	O View all recent events for any debugging
kubectl get eventssort-by=.metadata.creationTimestamp	Events sorted by time (great for finding root cause)
kubectl taint nodes <node> key=value:NoSchedule</node>	O Prevent pods from scheduling on specific node (node maintenance)

Command CPU/metrics	Purpose
kubectl top podcontainers	See CPU/memory usage per container in pod
kubectl top nodeuse-protocol-buffers	More accurate node metrics
kubectl getraw "/apis/metrics.k8s.io/v1beta1/nodes"	Raw metrics endpoint — useful for scripting

Command	Purpose
kubectl get sa	List service accounts in current namespace
kubectl describe sa <name></name>	Show tokens and bindings
kubectl get roles,rolebindings,clusterroles,clusterrolebindings	Full RBAC scope
kubectl create secret generic <name>from- literal=password=admin123</name>	Create a basic secret
`kubectl get secrets -o jsonpath="{.items[*].data}"	base64 decode secrets manually`

What is Ingress in Kubernetes?

An Ingress is an API object in Kubernetes that manages external access to services in a cluster, typically HTTP/HTTPS routes. It lets you define rules for routing traffic to backend services based on the request host or path—think of it like a smart router for your services.

What is an Ingress Controller?

An Ingress Controller is the actual implementation that reads the Ingress rules and configures a reverse proxy (like NGINX, Traefik, etc.) to enforce those rules.

Without an Ingress Controller, your Ingress rules do nothing.

How it works:

You create an Ingress resource with rules (e.g., route /api to Service A and /web to Service B).

The Ingress Controller sees those rules and configures its proxy (e.g., NGINX) to implement them.

Incoming requests hit the Ingress Controller's LoadBalancer IP.

The controller routes the request based on the Ingress rules.

Ingress controller installation:

kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controllerv1.9.1/deploy/static/provider/cloud/deploy.yaml

Component	Description
Ingress	Defines routing rules to expose internal services
Ingress Controller	Listens for Ingress resources and configures the proxy
NGINX/Traefik	Popular Ingress Controllers (implement the logic)
LoadBalancer/NodePort	Expose the Ingress Controller to the internet

Note: Ingress by itself cannot do anything. It relies on the presence of an **Ingress Controller**. Once an Ingress Controller is installed, it watches for Ingress resources and generates all the necessary configurations internally. The Ingress communicates with the Ingress Controller, which then routes traffic based on **host**, **path**, or **network-based** rules defined in the Ingress resource.

Relationship between Ingress and Ingress Controller

- •Ingress is just a set of rules that define how external HTTP/HTTPS traffic should be routed to services inside your Kubernetes cluster.
- •Ingress Controller is the actual implementation (like NGINX, Traefik, or AWS ALB) that reads those Ingress rules and configures a reverse proxy to route the traffic accordingly.

Marks How it works:

- 1. You define an **Ingress** resource with routing rules.
- 2. The **Ingress Controller** constantly watches the Kubernetes API for new or updated Ingress rules.
- 3.It **generates configuration** (e.g., for NGINX) based on those rules.
- 4.It handles incoming requests and routes them to the correct service.

```
# ingress.yaml
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
 name: hello-ingress
 annotations:
  nginx.ingress.kubernetes.io/rewrite-target: /
spec:
 rules:
 - host: hello.example.com
  http:
   paths:
   - path: /
    pathType: Prefix
    backend:
     service:
      name: hello-service
       port:
        number: 80
```

StatefullSet and Stateless vs DeamonSet:



Use When:

You need persistent storage, stable network identities, or ordered deployment (e.g., for databases like MongoDB, Cassandra, Kafka).

Features:

Persistent Volumes: Each pod gets its own storage (PVC) that persists across restarts.

Stable DNS names: Pods get predictable names like pod-0, pod-1 etc.

Ordered scaling: Pods are created/terminated in order.

Example Use Case:

Databases, message brokers, clustered applications.



b Use When:

You don't need to store data locally, and pods can be easily replaced. Most web apps, APIs, or frontend services fall here.

Features:

Pods are interchangeable (no identity).

No guaranteed order of deployment or termination.

Easy scaling and rolling updates.

Example Use Case:

Web servers, REST APIs, microservices.

Use When:

You want to run **one pod on every (or selected) node** in the cluster.

- Features:
- •Automatically deploys one pod per node.
- •Used for **cluster-wide services** like monitoring, logging, networking agents.
- •Pods automatically appear on **new nodes** added to the cluster.
- Example Use Case:
- •Fluentd/Logstash for logging.
- •Prometheus Node Exporter for monitoring.
- •CNI plugins, or security agents.

StatefulSet = For apps that need identity + persistent data.

♣ Deployment = For stateless apps that can be scaled easily.

DaemonSet = For running a pod on every node, usually for system-level tasks.

Feature	StatefulSet	Deployment (Stateless)	DaemonSet
Pod Identity	Sticky, stable (e.g., pod-0)	No identity (all equal)	One pod per node
Storage	Unique, persistent per pod	Shared or none	Optional, often none
Use Case	Databases, Kafka	Web apps, APIs	Monitoring, logging agents
Scaling	Ordered	Random	Based on node count
Pod Per Node	No	No	Yes

StatefulSet YAML (with PVC)

apiVersion: apps/v1 kind: StatefulSet metadata: name: my-db spec: serviceName: "my-db-headless" replicas: 3 selector: matchLabels: app: my-db template: metadata: labels: app: my-db spec: containers: - name: db image: mongo:5 ports: - containerPort: 27017 volumeMounts: - name: db-storage mountPath: /data/db volumeClaimTemplates: - metadata: name: db-storage spec: accessModes: ["ReadWriteOnce"] resources: requests:

storage: 1Gi

Stateless Deployment YAML:

apiVersion: apps/v1 kind: Deployment metadata: name: web-app spec: replicas: 3 selector: matchLabels: app: web template: metadata: labels: app: web spec: containers: - name: web image: nginx:latest ports:

- containerPort: 80

DaemonSet YAML:

apiVersion: apps/v1 kind: DaemonSet metadata: name: node-monitor spec: selector: matchLabels: name: node-monitor template: metadata: labels:

Master node







Worker

nodes

Through DaemonSet POD schedule

In all Nodes

containers:

spec:

name: node-exporter image: prom/node-exporter:latest ports:

- containerPort: 9100

name: node-monitor

A pod is created on every node automatically.

Useful for cluster-wide agents like metrics/log collection.

Ways to Schedule Pods on Selected or All Nodes:

1. DaemonSet

kind: DaemonSet
spec:
template:
spec:
nodeSelector:
node-role.kubernetes.io/worker: ""

Duse Case: Logging agents, monitoring exporters (e.g., Prometheus Node Exporter).

2. **nodeSelector** – Run pods on SELECTED nodes

Assign pods to specific nodes using **key-value labels**.

spec:
 nodeSelector:
 disktype: ssd

Use Case: Run a database only on nodes with SSDs.

3. nodeAffinity – More advanced node matching

★ Use Case: Better control and flexibility over scheduling compared to nodeSelector. spec:
 affinity:

nodeAffinity: requiredDuringSchedulingIgnoredDuringExecution:

nodeSelectorTerms:

- matchExpressions:

key: disktype operator: In values:

- ssd

4. **Taints & Tolerations** – Allow or prevent pods from running on nodes

> Taints: Mark a node as restricted.

> Tolerations: Allow pods to "tolerate" the taint and run there.

Taint on node: kubectl taint nodes node1 key=value:NoSchedule

Toleration in pod:

tolerations:
- key: "key"
operator: "Equal"
value: "value"
effect: "NoSchedule"

Use Case: Control workloads (e.g., only system pods on master nodes).



CoreDNS is the **default DNS server** used by Kubernetes (since v1.13+) for **cluster DNS resolution**. It helps pods and services **communicate using names** instead of IP addresses.

Term	Meaning
DNS	Domain Name System – converts names to IPs
CoreDNS	A DNS server written in Go and designed to be extensible
kube-dns	The older DNS service (replaced by CoreDNS)



What CoreDNS Does in Kubernetes

Listens for DNS queries from pods.

Resolves service names like my-svc.default.svc.cluster.local to a cluster IP.

Works with kubelet and kube-proxy for name-to-IP mapping.

Uses plugins (like kubernetes, forward, cache, etc.) for flexibility.

How resolution works:

CoreDNS receives the query: "What is the IP of my-svc.default.svc.cluster.local?" It then looks up the Kubernetes API to find the service my-svc in the default namespace.

If found, it returns the ClusterIP (e.g., 10.97.53.42) to the pod that requested it The pod uses this IP to connect to the service.

Corefile (CoreDNS Configuration):

```
.:53 {
  errors
  health
  kubernetes cluster.local in-addr.arpa ip6.arpa {
    pods insecure
    fallthrough in-addr.arpa ip6.arpa
 forward . /etc/resolv.conf
  cache 30
  loop
  reload
  loadbalance
```

1 Listens for DNS queries from Pods

Every Pod in Kubernetes has a file called /etc/resolv.conf that tells it where to send DNS queries.

This file usually contains:

nameserver 10.96.0.10 # IP of the CoreDNS service (cluster IP)

When a pod wants to communicate with another service (e.g., my-service.default.svc.cluster.local), it performs a DNS query. That DNS query is sent to CoreDNS, which is running as a service in the kube-system namespace.

2 Resolves service names like my-svc.default.svc.cluster.local to ClusterIP

Example Breakdown:

my-svc.default.svc.cluster.local

my-svc: the name of the service This enables service discovery using names instead of hardcoded IPs.

default: the namespace

svc: indicates it's a service

cluster.local: the cluster domain (default DNS suffix)

3 Works with kubelet and kube-proxy for name-to-IP mapping

Component	Role
kubelet	Configures the pod's networking and injects the correct DNS settings (like pointing to CoreDNS IP).
kube-proxy	Creates IP tables or IPVS rules so that the ClusterIP returned by CoreDNS routes traffic to one of the backend pods for that service.

4 Uses Plugins for Flexibility

CoreDNS is modular, and its behavior is defined in the Corefile (usually stored in the CoreDNS ConfigMap).

Common plugins and their roles:

```
apiVersion: v1
kind: ConfigMap
metadata:
 name: coredns
 namespace: kube-system
data:
 Corefile: |
  .:53 {
    errors
    health
    kubernetes cluster.local in-addr.arpa ip6.arpa {
      pods insecure
      fallthrough in-addr.arpa ip6.arpa
    forward . /etc/resolv.conf
    cache 30
    loop
    reload
    loadbalance
```



How a Pod Uses CoreDNS

Pod is created with /etc/resolv.conf pointing to CoreDNS.

When a pod tries to connect to my-service.default.svc.cluster.local, it sends a DNS query.

CoreDNS checks its internal rules and returns the correct IP.

The pod uses that IP to connect to the service.

- \$\left\tau \cdots: 153 means listen on port 53 for all domains.
- forward . /etc/resolv.conf = forward unknown queries (e.g., google.com) to the host's DNS.

Plugins for Flexibility:

Plugin	What it does	
kubernetes	Enables service and pod name resolution using the Kubernetes API.	
forward	Forwards unresolved DNS queries to an external DNS (like Google DNS or the node's /etc/resolv.conf).	
cache	Caches DNS responses to improve speed and reduce API load.	
loop	Detects and prevents infinite DNS query loops.	
health	Exposes an HTTP health endpoint for CoreDNS.	
reload	Auto-reloads CoreDNS if its config changes.	
loadbalance	Randomizes the list of A records returned (helps in load balancing across pods).	



A ServiceAccount (SA) is an identity used by pods to interact with the Kubernetes API.



This account is automatically mounted into the pod as a token in /var/run/secrets/kubernetes.io/serviceaccount/.



Pods often need to interact with the Kubernetes API (e.g., to read config maps, list pods, watch services). Instead of using user credentials (bad practice), they use service accounts.

ServiceAccount.yaml

apiVersion: v1

kind: ServiceAccount

metadata:

name: custom-sa

namespace: default

apiVersion: v1 kind: Pod

metadata:

name: mypod

#POD Attaching in Service Account

spec:

serviceAccountName: custom-sa

containers:

name: nginx image: nginx apiVersion: apps/v1 kind: Deployment metadata:

name: app-deploy

spec: replicas: 2 selector: matchLabels:

app: testapp template: metadata:

labels:

app: testapp

spec:

serviceAccountName: custom-sa

containers:
- name: app
image: nginx

How a Deployment Uses a Service Account

When you create a Deployment, each Pod it creates will use a Service Account to interact with the Kubernetes API.

By default, pods use the default service account in the same namespace, unless overridden.



How RBAC Protects Kubernetes Resources Using Service Accounts

RBAC controls what the service account is allowed to do.

If your app inside the pod tries to:

- •List all pods
- •Get secrets
- •Create deployments
- Access nodes



It will only succeed if RBAC allows it.



What is RBAC (Role-Based Access Control)?

RBAC lets you define what actions (verbs) a user, group, or service account can perform on which resources in Kubernetes. RBAC works using 4 main components:

Component	What it does	
Role	Defines permissions within a namespace	
ClusterRole	Defines cluster-wide permissions (or reusable across namespaces)	
RoleBinding	Grants a Role to a user/SA in a namespace	
ClusterRoleBinding	Grants a ClusterRole to a user/SA across the entire cluster	

Create Role:

apiVersion: rbac.authorization.k8s.io/v1

kind: Role metadata:

namespace: default

name: pod-reader

rules:

- apiGroups: [""]

resources: ["pods"]

verbs: ["get", "list", "watch"]

Bind Role to a ServiceAccount:

apiVersion: rbac.authorization.k8s.io/v1

kind: RoleBinding

metadata:

name: read-pods-binding

namespace: default

subjects:

- kind: ServiceAccount

name: custom-sa namespace: default

roleRef: kind: Role

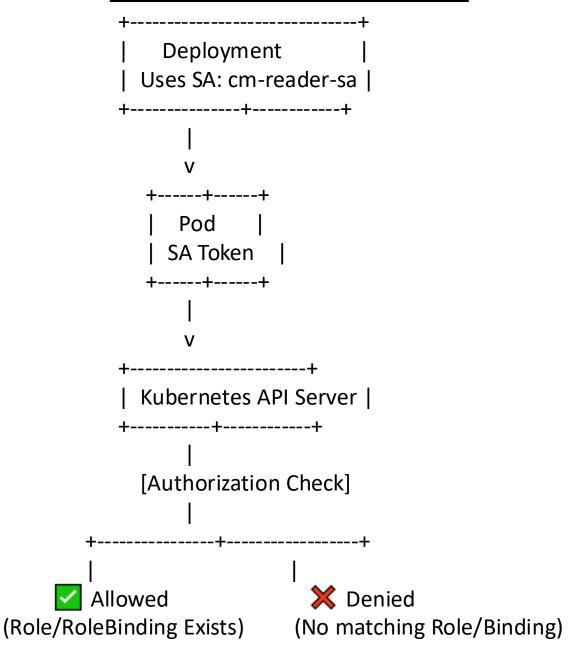
name: pod-reader

apiGroup: rbac.authorization.k8s.io

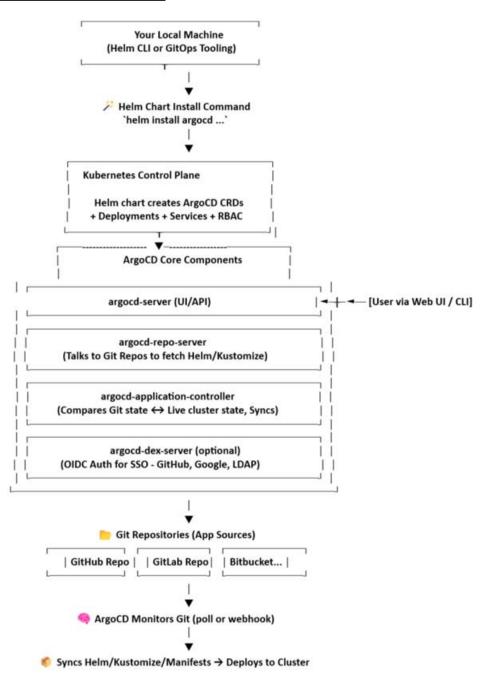
	Now custom-sa	can only read	pods in the	default
na	mespace.			

Task	Use	
Give pod permission to list secrets	Create SA + Role + RoleBinding	
Let Jenkins/ArgoCD deploy to namespaces	Use ClusterRole + ClusterRoleBinding	
Protect critical namespaces (e.g., kube-system)	Use RBAC to restrict access	

Service aaccount Deployment RBAC



Argocd Deployment





Helm is a package manager for Kubernetes that simplifies deployment and management of applications using Helm charts. It

allows you to:

- Deploy applications with a single command
- Use templates to manage configurations
- Version control your deployments
- Easily upgrade/rollback applications



Helm Charts – Pre-packaged application definitions (like a blueprint).

Values.yaml – Configuration file for customizing deployments.

Templates/ – Kubernetes manifests with dynamic placeholders ({{ }}) for customization.

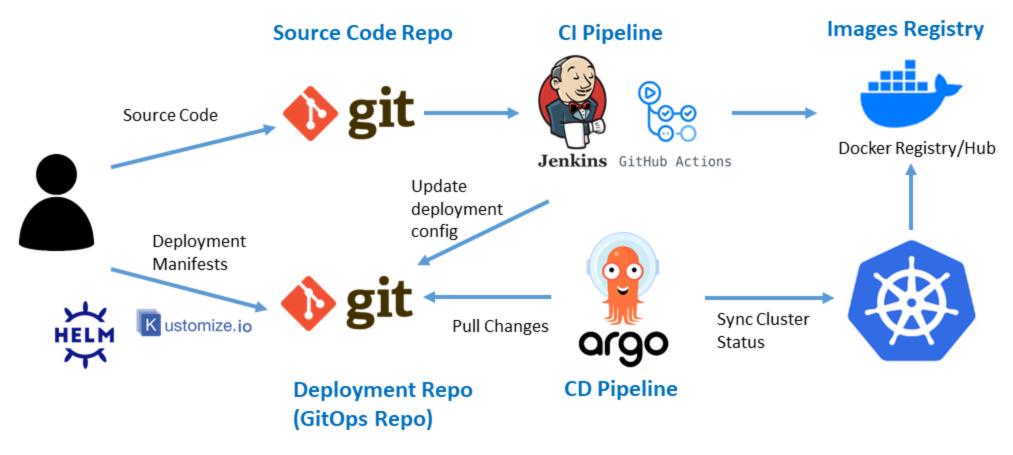
Releases – A deployed instance of a Helm chart.

Repositories – Where Helm charts are stored and shared.

My manifest:

```
apiVersion: apps/v1
kind: Deployment
metadata:
name: {{ .Release.Name }}-app
spec:
 replicas: {{ .Values.replicaCount }}
 selector:
  matchLabels:
   app: {{ .Release.Name }}-app
template:
  metadata:
   labels:
    app: {{ .Release.Name }}-app
  spec:
   containers:
    - name: {{ .Chart.Name }}
     image: "{{ .Values.image.repository }}:{{ .Values.image.tag }}
     ports:
      - containerPort: 8080
```

Application Deployment Process:



Thank You....!

B SURYA PRAKASH REDDY