CAREER BYTE CODE

REALTIME PROJECTS PLATFORM

91 COUNTRIES 241k Learners

+32 471 40 89 08

CAREERBYTECODE.SUBSTACK.COM

RealTime Handson

careerbytecode

100 Kubernetes REALTIME SCENARIO BASED QUESTIONS

techopsbysonali

careerbytecode.substack.com

1. You need to gracefully terminate pods to avoid disrupting ongoing processes during a deployment. How would you configure Kubernetes to ensure graceful pod termination? What role do preStop hooks and termination grace periods play in this process?

2. Your team has implemented a custom resource definition (CRD), but it isn't behaving as expected. How would you debug and resolve issues with a CRD?
What tools (e.g., kubectl describe, logs, webhook configurations) can help you identify the problem?

3. The Kubernetes API server becomes unresponsive in your cluster.
How would you diagnose and restore API server functionality?
What precautions can you take to minimize the impact of API server downtime?

4. Dealing with Node Affinity and Anti-Affinity Conflicts You've configured node affinity and anti-affinity rules for a workload, but pods are not being scheduled as expected. How would you troubleshoot these scheduling issues? What are best practices for using affinity and anti-affinity to manage workload placement?

5. Application pods in your cluster are generating large log files, causing storage issues. How would you manage log rotation and retention for pods?
What Kubernetes features or external tools can assist in log management?

6. Your team is migrating to Helm for application deployment, but some charts require custom values. How would you manage and deploy custom values for Helm charts? What strategies can you use to handle versioning and rollbacks for Helm releases?

7. A StatefulSet for your database application is failing to create pods. How would you debug why the StatefulSet is failing? What are the specific considerations for storage and network configurations in StatefulSets?

8. Your cluster uses a mix of third-party and custom container images. How would you ensure that only secure and compliant images are used in the cluster? What tools (e.g., image scanners, policies) can assist with this process?

9. You deployed a Terraform module that creates a Kubernetes cluster along with necessary resources for an application. However, one of your Pods is stuck in a "CrashLoopBackOff" state. How would you troubleshoot this issue using Terraform and Kubernetes?

10. Your organization wants to ensure no single team can overuse cluster resources. How would you implement resource quotas and limits using Terraform in a multi-namespace Kubernetes cluster?

11. You have an e-commerce application running on Kubernetes, managed via Terraform. During a flash sale, the application faces high traffic, and the Pods are unable to handle the load. How would you configure horizontal scaling to handle this spike?

12. A deployment managed via Terraform fails, and the new version of your application is not stable. The rollout is stuck, impacting the service. How would you identify and mitigate the issue while ensuring minimal downtime?

13. You've deployed a microservices-based architecture using Terraform, but one service cannot communicate with another. How would you debug and resolve the networking issue?

14. A service configured with a LoadBalancer type is not accessible from outside the cluster. How would you debug and resolve this issue?
What are the key components and configurations to check?

15. You've enabled the Cluster Autoscaler in your cloud provider's Kubernetes cluster, but it isn't adding nodes even when resources are insufficient.
How would you debug why the Cluster Autoscaler isn't working?
What configurations and prerequisites are necessary for the autoscaler to function correctly?

16. You observe that certain nodes in your cluster are underutilized, leading to resource wastage. How would you identify and address the issue of underutilized nodes?What strategies, like overprovisioning or pod scheduling, can optimize resource usage?

17. Your team wants to test new application versions with a small percentage of users before a full rollout. How would you implement a canary deployment in Kubernetes?What tools (e.g., Argo Rollouts, Flagger) can assist in this process, and how do they work?

18. You need to manage application configurations that change frequently without redeploying the application. How would you use ConfigMaps and Secrets to achieve this?What are the differences between these resources, and how do you securely manage them?

19. A production application deployed on Kubernetes is experiencing intermittent failures. How would you debug a live application without disrupting other users?Which Kubernetes tools (e.g., kubectl debug, logs, metrics) would you leverage?

20. Your organization requires a highly available application that spans multiple regions. How would you design and deploy a Kubernetes-based solution for cross-region high availability? What challenges would you face, and how would you mitigate them?

21. You discover that Kubernetes certificates in your cluster are about to expire.
How would you identify and renew expiring certificates?What tools or steps would you use to prevent this from happening in the future?

22. Your organization mandates encrypted communication between pods.
How would you implement pod-to-pod encryption in a Kubernetes cluster?
What tools, such as mTLS or service meshes, would you use, and why?

23. Your cluster's scheduler is taking too long to assign pods to nodes.
How would you debug and optimize scheduler performance?What Kubernetes features (e.g., PriorityClass, scheduler configuration) can improve scheduling efficiency?

24. Your application is set up with a Horizontal Pod Autoscaler (HPA), but scaling is not happening even under high load. How would you troubleshoot why the HPA is not scaling the pods?What are the prerequisites for HPA to function properly?

25. You are deploying a stateful application, such as a database, on Kubernetes.
What are the key differences between StatefulSets and Deployments, and why would you choose one over the other?How do you handle scaling and backups for stateful workloads?

26. Your cluster's API server is responding slowly, impacting other components.
How would you diagnose and resolve API server performance bottlenecks?
What are the common causes of high API server latency?

27. A pod is stuck waiting for its Persistent Volume Claim (PVC) to be bound.
How do you debug and resolve PVC binding issues?What are the key considerations when provisioning storage dynamically in Kubernetes?

28. You need to configure a Kubernetes cluster for multi-tenancy to isolate workloads from different teams. How would you implement multi-tenancy in Kubernetes?What tools or features would you use to enforce resource isolation and security?

29. Your application pods are taking too long to start.What could be causing the slow startup, and how would you debug the issue?How do liveness and readiness probes impact pod startup?

30. Pods in your cluster are unable to resolve external domain names.
How would you debug and resolve DNS resolution failures in Kubernetes?
What are the key components involved in DNS resolution in a Kubernetes cluster?

31. A namespace in your cluster has reached its resource quota, and new pods can't be scheduled.
How would you diagnose and resolve the issue?

32. Your security team mandates that only images from a trusted private registry can be used in your Kubernetes cluster. How would you enforce this policy in your cluster? What Kubernetes features or tools can be used to achieve this?

33. Your team decides to implement a service mesh for better observability, security, and traffic control between microservices. How would you introduce a service mesh like Istio or Linkerd into your Kubernetes environment? What challenges would you expect during implementation, and how would you address them?

34. You are tasked with upgrading your Kubernetes cluster to a newer version.
What challenges might you face during a cluster upgrade?
How would you plan and execute a seamless upgrade without impacting running applications?

35. Your organization requires stricter security controls for pod deployments.
How would you enforce pod security policies to restrict permissions, such as root access or privileged containers?
What alternatives to PSPs would you consider since they are deprecated?

36. You need to monitor application performance and resource usage in your cluster.Which tools would you use to monitor Kubernetes clusters effectively?
How would you implement monitoring for resource usage, application logs, and network traffic?

37. A critical application is deployed on your Kubernetes cluster, and you are required to implement disaster recovery. How would you design a backup and recovery solution for the application and cluster?
What tools or strategies would you use to back up etcd, persistent data, and application configurations?

38. You need to perform maintenance on a node in your cluster, but it runs critical application pods.
How would you drain the node without causing downtime for the applications? What would you do if some pods fail to evict during the drain?

39. You have applications running in two separate Kubernetes clusters, and they need to communicate securely. How would you establish secure communication between applications across clusters? What tools or practices would you use for cross-cluster networking?

40. Your team demands zero downtime during deployments.
How would you configure your deployment strategy to achieve zero downtime?
Compare and contrast the use of RollingUpdate and Blue/Green Deployments.

41. You need to expose an application securely over HTTPS using an ingress controller. How would you configure an ingress resource with SSL termination?
What considerations would you make for managing SSL certificates in Kubernetes?

42. You have been asked to review the security posture of your Kubernetes cluster.What steps would you take to harden the cluster?
How would you ensure RBAC policies, network policies, and secure communication between components are implemented?

43. A pod is failing to start because it can't pull the required container image.
How would you debug the image pull issue?What are the possible reasons for such failures, and how can you mitigate them?

44. Your organization requires a highly available application that spans multiple regions.How would you design and deploy a Kubernetes-based solution for cross-region high availability?What challenges would you face, and how would you mitigate them?

45. You discover that Kubernetes certificates in your cluster are about to expire.How would you identify and renew expiring certificates?What tools or steps would you use to prevent this from happening in the future?

46. Your organization mandates encrypted communication between pods.How would you implement pod-to-pod encryption in a Kubernetes cluster?What tools, such as mTLS or service meshes, would you use, and why?

47. Your cluster's scheduler is taking too long to assign pods to nodes.How would you debug and optimize scheduler performance?What Kubernetes features (e.g., `PriorityClass`, `scheduler configuration`) can improve scheduling efficiency?

48. You need to gracefully terminate pods to avoid disrupting ongoing processes during a deployment.How would you configure Kubernetes to ensure graceful pod termination? What role do preStop hooks and termination grace periods play in this process?

49. Your team has implemented a custom resource definition (CRD), but it isn't behaving as expected.How would you debug and resolve issues with a CRD?What tools (e.g., `kubectl describe`, logs, webhook configurations) can help you identify the problem?

50. The Kubernetes API server becomes unresponsive in your cluster. How would you diagnose and restore API server functionality? What precautions can you take to minimize the impact of API server downtime?

51. Your Kubernetes cluster experiences high CPU and memory consumption, affecting application performance.How would you identify which pods or nodes are consuming excessive resources? What actions would you take to resolve the issue and optimize resource usage?

52. A critical pod in production is stuck in a `CrashLoopBackOff` state.How would you troubleshoot and resolve this issue? What are common causes of `CrashLoopBackOff`, and how can they be prevented?

53. Some pods cannot communicate with each other, but others in the same namespace work fine. How would you debug a potential network partition in Kubernetes? What role do network policies, CNI plugins, and kube-proxy play in troubleshooting?

54. You need to rotate API keys stored in Kubernetes Secrets without causing application downtime. How would you implement a secret rotation strategy? What Kubernetes features or tools (e.g., CSI, external secrets) would you leverage?

55. Your company is looking to reduce Kubernetes infrastructure costs while maintaining performance. What strategies would you use to optimize cluster costs? How would you analyze and scale workloads efficiently?

56. A pod using a Persistent Volume (PV) gets rescheduled to a new node, but the volume does not reattach.How would you troubleshoot and resolve this issue? What considerations should be taken when designing PVs for high availability?

57. A microservice deployed in Kubernetes experiences intermittent timeouts when calling another service.How would you debug service timeouts and latency issues?What Kubernetes features (e.g., readiness probes, Istio, circuit breakers) can help prevent such problems?

58. Some pods in your cluster are getting terminated with `OOMKilled` errors. How would you identify and fix Out of Memory (OOM) issues in Kubernetes? What are the best practices for setting resource requests and limits?

59. One of your worker nodes fails unexpectedly, causing application downtime. How does Kubernetes handle node failures, and what steps can you take to minimize downtime? What features like node taints, eviction policies, and pod disruption budgets can help?

60. Some pods in your cluster take significantly longer to start than expected.How would you investigate and resolve slow startup times? What impact do init containers, probes, and resource constraints have on pod startup?

61. Your team is hitting Docker Hub's image pull rate limits, causing deployment failures. How would you mitigate this issue in a Kubernetes cluster? What alternatives to Docker Hub can be used for image distribution?

62. A newly deployed policy via an Admission Controller is blocking deployments unexpectedly.How would you troubleshoot and debug this issue? What tools and logs would you check for more insights?

63. Your Kubernetes API server is under high load, slowing down requests and causing timeouts. How would you diagnose and improve API server performance in a large cluster? What role do etcd tuning, API caching, and request limits play in optimizing performance?

64. Your team needs to migrate a database running as a StatefulSet to a new Kubernetes cluster. What challenges would you face when migrating StatefulSets? How would you ensure data consistency during the migration?

65. Pods in your cluster keep getting evicted, even though there seem to be enough resources available. How would you debug unexpected pod evictions? What Kubernetes events and logs would you analyze?

66. You need to enforce authentication between services running inside a Kubernetes cluster. How would you implement mutual TLS (mTLS) between services? What role do service meshes like Istio play in securing service-to-service communication?

67. Your company wants to deploy a multi-region active-passive Kubernetes architecture for disaster recovery. How would you design and implement automatic failover between regions? What tools and networking strategies would you use for seamless migration?

68. A scheduled Kubernetes CronJob is not running at the expected time. How would you troubleshoot and fix delayed or missing job executions? What factors can cause Kubernetes Jobs to fail unexpectedly?

69. Your team needs to upgrade Kubernetes from v1.24 to v1.26 with zero downtime.How would you plan and execute a Kubernetes upgrade while minimizing disruption?What considerations should be made for API deprecations and backward compatibility?

70. Your security team wants to restrict network traffic between different namespaces. How would you define and apply Kubernetes Network Policies? How can you test if network policies are working as expected?

71. A new deployment is scheduled, but the pods remain in a `Pending` state. How would you diagnose and resolve the issue? What are the key factors (e.g., node availability, resource requests, taints/tolerations) that can cause pods to remain pending?

72. Pods are unable to resolve domain names for internal or external services. How would you troubleshoot DNS issues in a Kubernetes cluster? What are the key components of Kubernetes DNS, and how can you debug them?

73. You notice that pods are underperforming despite having sufficient CPU requests and limits set. How would you identify and address CPU throttling? What are best practices for setting CPU requests and limits to avoid throttling?

74. Some nodes in your cluster are consistently running out of memory and CPU, causing application crashes. How would you diagnose and prevent resource starvation? What role do Kubernetes resource quotas, limits, and vertical pod autoscalers play?

75. Your entire Kubernetes cluster suddenly becomes unresponsive. What steps would you take to diagnose and recover from a full-cluster outage? What are some preventive measures to ensure high availability?

76. Your team wants to implement a blue-green deployment strategy for zero-downtime updates. How would you set up blue-green deployments in Kubernetes? What are the advantages and potential risks of this strategy?

77. A pod requesting a Persistent Volume Claim (PVC) is failing to get a bound volume. How would you troubleshoot and resolve PVC binding issues? What are common reasons for PVC binding failures?

78. Your team wants to combine horizontal and vertical pod autoscaling for an application. How would you configure both HPA and VPA to work effectively? What conflicts might arise, and how can they be mitigated?

79. Some pods in your cluster are restarting frequently, causing instability. How would you investigate and fix frequent pod restarts? What are common causes of pod flapping, and how can they be prevented?

80. Your application is experiencing high disk latency while reading and writing to a persistent volume.How would you troubleshoot and optimize storage performance in Kubernetes? What factors (e.g., storage class, volume type, provisioner) impact performance?

81. A newly deployed application is failing to start with an `ImagePullBackOff` error.
How would you troubleshoot this issue? What are some ways to prevent image pull failures in production?

82. Your cluster supports multiple teams, but some teams report issues with access controls and resource quotas.How would you enforce isolation and access control for multiple teams in a shared Kubernetes cluster? What Kubernetes features (e.g., RBAC, namespaces, network policies) can help manage multi-tenancy effectively?

83. Your application is exposed through an API gateway, but some requests are failing intermittently. How would you debug API gateway failures in Kubernetes?What strategies (e.g., rate limiting, retries, health checks) can improve API reliability?

84. An application is still using an old configuration even after updating the ConfigMap. Why might a pod not reflect the latest ConfigMap changes? How can you ensure that applications always use the latest configuration updates?

85. A developer reports they cannot create deployments, even though they have `edit` access in a namespace.How would you debug and fix RBAC-related access issues? What are best practices for managing RBAC policies in a Kubernetes cluster?

86. A database running as a StatefulSet in Kubernetes is experiencing data inconsistencies across replicas.How would you debug and resolve replication issues in a StatefulSet? What are best practices for running stateful applications in Kubernetes?

87. Your team needs to centralize and analyze logs from all applications running in Kubernetes.
How would you design and implement a log aggregation solution? What logging tools (e.g., Fluentd, Loki, Elasticsearch) are best suited for large-scale Kubernetes environments?

88. Your security team requires strict controls over the security context of pods.How would you enforce security policies for pods? What tools (e.g., PodSecurityAdmission, OPA Gatekeeper) can help implement security best practices?

89. Your team is implementing a backup strategy for Kubernetes applications and PersistentVolumes. How would you design a backup and restore process for Kubernetes workloads? What tools (e.g., Velero, Stash) can help with Kubernetes backup and disaster recovery?

90. A DaemonSet is expected to run on all nodes, but some nodes are missing pods.How would you troubleshoot why certain nodes do not have DaemonSet pods running? What are common issues that prevent DaemonSet scheduling?

91. Your company wants to deploy Kubernetes workloads across multiple cloud providers (AWS, GCP, Azure). How would you design and manage a multi-cloud Kubernetes architecture? What are the challenges of running Kubernetes clusters across multiple clouds, and how can they be addressed?

92. Your CI/CD pipeline deploys container images to Kubernetes, but you need to enforce security scans.How would you integrate image scanning into the CI/CD workflow? What tools (e.g., Trivy, Clair, Aqua Security) can help enforce security best practices?

93. A recent canary deployment caused service downtime instead of a gradual rollout. How would you troubleshoot and fix issues with a failed canary release?What tools (e.g., Argo Rollouts, Flagger) can help implement safe canary deployments?

94. Your cloud provider requires Kubernetes worker nodes to be upgraded. How would you perform a rolling upgrade of worker nodes without causing downtime? What are the risks involved, and how can they be mitigated?

95. You need to enforce HTTPS and apply WAF rules for applications exposed via an Ingress controller. How would you secure traffic handled by the Kubernetes Ingress controller? What are the best practices for TLS termination, authentication, and DDoS protection?

96. Your application uses an InitContainer to set up configurations before the main container starts. However, pods are failing because the InitContainer is stuck in `CrashLoopBackOff`. How would you diagnose and fix this issue? What are common reasons for InitContainer failures?

97. A Kubernetes Job is scheduled but never completes, and the pod remains in a `Running` state indefinitely. How would you troubleshoot this issue? What are best practices for defining Job configurations to prevent such problems?

98. Your Kubernetes API server suddenly stops responding to requests, affecting cluster operations. What steps would you take to investigate and restore API server functionality? How can you ensure high availability for the Kubernetes API server?

99. Your application is experiencing high response times, and resource utilization is inconsistent across nodes. How would you identify and resolve performance bottlenecks? What Kubernetes monitoring tools (e.g., Prometheus, Grafana) would you use for performance analysis?

100. A microservices-based application deployed in Kubernetes is experiencing intermittent network delays. How would you diagnose and fix network latency issues? What role do Kubernetes CNI plugins and Service Meshes (e.g., Istio, Linkerd) play in optimizing network performance?

CAREER BYTE CODE
REALTIME PROJECTS PLATFORM

91 COUNTRIES 241k Learners
+32 471 40 89 08
CAREERBYTECODE.SUBSTACK.COM

→ TRAININGS

WE ARE DIFFERENT

At CareerByteCode, we redefine training by focusing on real-world, hands-on experience. Unlike traditional learning methods, we provide step-by-step implementation guides, 500+ real-time use cases, and industry-relevant projects across cutting-edge technologies like AWS, Azure, GCP, DevOps, AI, FullStack Development and more.

Our approach goes beyond theoretical knowledge—we offer expert mentorship, helping learners understand how to study effectively, close career gaps, and

16+
Years of operations

91+
Countries worldwide

241 K Happy clients

Our Usecases Platform
https://careerbytecode.substack.com

Our WebShop

CareerByteCode
Learning Made simple

CareerByteCode
All in One Platform

STAY IN TOUCH WITH US!

Website

Our WebShop https://careerbytecode.shop
Our Usecases Platform https://careerbytecode.substack.com

E-mail
careerbytec@gmail.com

Social Media
@careerbytecode

Phone
+32 471 40 8908

HQ address
Belgium, Europe

CAREER BYTE CODE

REALTIME PROJECTS PLATFORM

91 COUNTRIES 241k Learners

+32 471 40 89 08

CAREERBYTECODE.SUBSTACK.COM

For any RealTime Handson Projects

And for more tips like this

Follow

Like & ReShare

@techopsbysonali